


Session 9: Simulation Modeling I (Solutions Only)


Q1. Sampling from Common Distributions

a) Generate 10000 samples of a uniform distribution between 1 and 3 and plot the histogram, as well as the empirical CDF. Calculate the mean and standard deviation of the samples, as well as the proportion between 2 and 2.5 (inclusive).

```
[10]: from scipy.stats import uniform
      s=pd.Series(uniform(1,2).rvs(10000))
      s.hist()
```



```
[11]: s.hist(bins=100,density=True,cumulative=True)
```



```
[12]: s.mean()
```

```
1.9973519125605361
```

```
[13]: s.std()
```

```
0.5797250400567447
```

```
[14]: ((s>=2) & (s<=2.5)).mean()
```

```
0.2421
```

Q1-b: Generate 100 samples of a binomial distribution with $n = 10$ and $p = 0.3$. Calculate the mean and standard deviation of the sample and compare with what it should be from the formula. Plot a histogram with 50 bins. Repeat with 10,000 samples.


```
[15]: from scipy.stats import binom
import numpy as np
n,p=10,0.3
dist=binom(n,p)
s=pd.Series(dist.rvs(10))
print('Sample mean:',s.mean())
print('Sample std:',s.std())
print('Theoretical mean:',n*p)
print('Theoretical std:',np.sqrt(n*p*(1-p)))
s.hist(bins=50)
```

Sample mean: 2.9

Sample std: 1.449137674618944

Theoretical mean: 3.0

Theoretical std: 1.4491376746189437




```
[16]: s=pd.Series(dist.rvs(10000))
print('Sample mean:',s.mean())
print('Sample std:',s.std())
print('Theoretical mean:',n*p)
print('Theoretical std:',np.sqrt(n*p*(1-p)))
s.hist(bins=50)
```

Sample mean: 2.999

Sample std: 1.455819680143839

Theoretical mean: 3.0

Theoretical std: 1.4491376746189437



Q2. Sampling from Mixture of Distributions

You are tasked with forecasting demand for a new product. Based on past data and your knowledge of the product, you estimate that the product quality will be Amazing with probability 0.1, Mediocre with probability 0.5, and Terrible with probability 0.4. You model the demand as normally distributed, with mean and standard deviation depending on the product quality as follows.

Prod. Quality:	Amazing	Mediocre	Terrible
μ	10000	5000	1000
σ	2000	1000	300


Create a Series called “forecast” with 10,000 samples of the demand forecast, and compute the mean and standard deviation of the samples, as well as the probability that demand is more than 6000. Finally, plot a histogram of the samples with 100 bins, as well as the empirical CDF.

```
[22]: from scipy.stats import binom
import numpy as np
data=[]
for i in range(10000):
    quality=np.random.choice(['A','M','T'],p=[0.1,0.5,0.4])
    if quality=='A':
        sample=norm(10000,2000).rvs()
    elif quality=='M':
        sample=norm(5000,1000).rvs()
    else:
        sample=norm(1000,300).rvs()
    data.append(sample)
forecast=pd.Series(data)
print('Sample mean:',forecast.mean())
print('Sample standard deviation:',forecast.std())
print('Probability demand more than 6000:',(forecast>6000).mean())
forecast.hist(bins=100)
```


Sample mean: 3947.671852127712

Sample standard deviation: 2965.1276604886903

Probability demand more than 6000: 0.1798



```
[23]: forecast.hist(bins=100,density=True,cumulative=True)
```



Q3. Sampling from More Complex Distributions

Nadeem is a car salesperson who faces the following incentive scheme at the dealership where he works. For each month, there is a “target profit” that the dealership sets for the month. If he makes more profit for the dealership that month than the target, then he receives a 20% bonus on the amount of profit over the target. However, if he does not meet the target, he receives no bonus. For example, if the target is 80,000 and he makes 100,000 of profit, then he receives a 4,000 bonus that month. However, if he makes 70,000, then he receives zero bonus that month. Nadeem would like to understand the distribution of his monthly bonus.

Nadeem estimates that the number of cars he sells is binomial distributed with $n = 200$ and $p = 0.2$. On every car he sells, the amount of profit he makes for the dealership is normally distributed with $\mu = 3000$ and $\sigma = 1000$, and the profit from each car is independent of another.


Create a Series called “monthlyBonus” with 10,000 samples of his monthly bonus. Compute the mean, the standard deviation, the probability the bonus is less than 5000, and plot a histogram with 50 bins as well as the empirical CDF.

```
[26]: import numpy as np
import pandas as pd
from scipy.stats import binom, norm
data=[]
target=80000
for i in range(10000):
    cars=binom(200,0.2).rvs()
    sales=np.sum(norm(3000,1000).rvs(size=cars))
    if sales>target:
        bonus=(sales-target)*0.2
    else:
        bonus=0
    data.append(bonus)
monthlyBonus=pd.Series(data)
print('Mean is',monthlyBonus.mean())
print('Standard deviation is',monthlyBonus.std())
print('Probability less than 5000 is',(monthlyBonus<5000).mean())
monthlyBonus.hist(bins=50)
```


Mean is 8003.109754601626

Standard deviation is 3595.790236751844

Probability less than 5000 is 0.2081



```
[27]: monthlyBonus.hist(bins=50,density=True,cumulative=True)
```



(Bonus) This question asks you to illustrate the Central Limit Theorem (CLT) by example. Consider the following distribution,

$$X = \begin{cases} 3 & \text{with probability 0.4,} \\ 5 & \text{with probability 0.1,} \\ 9 & \text{with probability 0.5.} \end{cases}$$


Define Y_n to be the sum of n independent random variables with the above distribution. Create 10000 samples of Y_1 , Y_5 , Y_{30} , Y_{100} , and Y_{1000} and plot their histograms (with 30 bins). You should be able to see the histograms converging to a Bell curve as n increases.

The Central Limit Theorem (CLT) says that this phenomenon always happens, regardless of the distribution of X . Moreover, it still takes place even if the n independent random variables do not have the same distribution, as long as each term in the sum is “small” relative to the whole. (For a precise mathematical formulation of the CLT in the case with non-identical random variables, search for Lyapunov or Lindeberg CLT on Wikipedia.)


```
[28]: import numpy as np
def simulateCLT(n,T=10000):
    x=[3,5,9]
    p=[0.4,0.1,0.5]
    data=[]
    for i in range(T):
        value=np.sum(np.random.choice(x,p=p,size=n))
        data.append(value)
    samples=pd.Series(data)
```

```
return samples
```


```
simulateCLT(1).hist(bins=30)
```




```
[29]: simulateCLT(5).hist(bins=30)
```



```
[30]: simulateCLT(20).hist(bins=30)
```



```
[31]: simulateCLT(100).hist(bins=30)
```



```
[32]: simulateCLT(1000).hist(bins=30)
```

