

DSO530 Statistical Learning Methods

Lecture 2a: Review and Simple Linear Regression

Dr. Xin Tong

Department of Data Sciences and Operations

Marshall School of Business

University of Southern California

Review: random variables and probability

- Suppose X is a standard normal random variable, i.e., $X \sim N(0, 1)$.
 - i). What is $P(X = 0)$?
 - ii). Can we say that $P(X = 0) > P(X = 1)$?
 - iii). For some small $\epsilon > 0$, is it true

$$P(X \in (-\epsilon, \epsilon)) > P(X \in (1 - \epsilon, 1 + \epsilon))?$$

- $Z \sim \text{Uniform}(0, 1)$. How much is $P(0.3 \leq Z \leq 0.7)$?
- We can visualize disjoint events A and B in a Venn Diagram. Is it clear how to visualize independent events?
- Are disjoint events independent?

Review: basics in Statistics and Lecture 1

- Can you name the difference between **supervised learning** and **unsupervised learning**?
- Define **population**, **sample**, **parameter**, **statistics** in the example:
“A politician selects a random sample of 200 working U.S. women who are 16 to 24 year old. Of the women in the sample, 5% are being paid minimum wage or less”.
- How to interpret **confidence intervals** (CIs)?
- Other things being equal, is 99% CI wider than 95%? Why we don't talk about 100% CI?

Review: basic concepts in Statistics

- What is **Simpson's Paradox**?
- Check out the hospital example in the week one technical notes and google the Berkeley graduate admission case.
- Simpson's Paradox vividly illustrates why business analytics must not be viewed as a purely technical subject appropriate for mechanization or automation.
- For further reading (3 cases):

[https://www.statslife.org.uk/the-statistics-dictionary/
2012-simpson-s-paradox-a-cautionary-tale-in-advanced-analytics](https://www.statslife.org.uk/the-statistics-dictionary/2012-simpson-s-paradox-a-cautionary-tale-in-advanced-analytics)

On more thing: about hypothesis test

- Hypothesis test: **null hypothesis** H_0 v.s. **alternative hypothesis** H_a
- Null hypothesis is usually the status quo, old technology, ineffective new drugs. . . .
- What is a **p-value**? the probability of obtaining a result equal to or “more extreme” than what was actually observed, when the null hypothesis is true
- α is a user specified value called *level of significance*
- We should design a statistical test such that $P_{H_0}(\text{reject } H_0) \leq \alpha$
- Reject the null hypothesis if p -value is smaller than or equal to α

On more thing: about hypothesis test

- Hypothesis test: **null hypothesis** H_0 v.s. **alternative hypothesis** H_a
- Null hypothesis is usually the status quo, old technology, ineffective new drugs. . . .
- What is a **p-value**? the probability of obtaining a result equal to or “more extreme” than what was actually observed, when the null hypothesis is true
- α is a user specified value called *level of significance*
- We should design a statistical test such that $P_{H_0}(\text{reject } H_0) \leq \alpha$
- Reject the null hypothesis if p -value is smaller than or equal to α

Linear regression with a single variable

- The package **scikit-learn** is a widely used Python library for machine learning, built on top of NumPy and some other packages.
- There are five basic steps when you are implementing linear regression:
 - Import the packages and classes you need.
 - Provide data to work with and do appropriate transformations if necessary.
 - Create a regression model and fit it with existing data.
 - Check the results of model fitting to know whether the model is satisfactory.
 - Apply the model for predictions.

Import the packages and read in data

```
: import numpy as np
  from sklearn.linear_model import LinearRegression
```

Figure 1: import NumPy and LinearRegression

```
import pandas as pd
import matplotlib.pyplot as plt
```

```
height = pd.read_csv("data/galton.csv");height.head()
```

	Unnamed: 0	child	parent
0	1	66.435917	70.851069
1	2	65.943364	69.858889
2	3	64.278858	65.278141
3	4	63.851914	64.032631
4	5	63.192294	63.418992

Figure 2: Do you see any undesirable elements in this DataFrame?

Yes, it is about the column 0!

```
: height = pd.read_csv("data/galton.csv", index_col=0); display(height.head())  
#can also try print(height.head())
```

	child	parent
1	66.435917	70.851069
2	65.943364	69.858889
3	64.278858	65.278141
4	63.851914	64.032631
5	63.192294	63.418992

Figure 3: This time it looks right

```
: height.info() # the .info() give you some information about the DataFrame height  
  
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 928 entries, 1 to 928  
Data columns (total 2 columns):  
child      928 non-null float64  
parent     928 non-null float64  
dtypes: float64(2)  
memory usage: 21.8 KB
```

Some summary statistics about the variables in the dataframe

```
height.describe() # columnwise summary statistics |
```

	child	parent
count	928.000000	928.000000
mean	68.086288	68.299524
std	1.527244	1.857460
min	63.192294	63.418992
25%	67.027340	67.233314
50%	68.115752	68.336807
75%	69.077425	69.485036
max	72.185332	73.355968

Figure 4

- What is the unit of measurements?
- What are the medians of the child and parent heights?

Visual inspection

```
plt.figure(); plt.scatter(height['parent'], height['child']); plt.show()
```




Figure 5: Scatter plot

- What information can be tell by looking at this scatter plot?

The most naive prediction

- Take parent's height as the predicted value for child's height: $\hat{y} = x$

```
: plt.figure(); plt.scatter(height['parent'], height['child'])  
plt.xlabel('parent height'); plt.ylabel('child height'); plt.title('child vs. parent heights')  
lineStart = height['parent'].min(); lineEnd = height['parent'].max()  
plt.plot([lineStart, lineEnd], [lineStart, lineEnd], color = 'r'); plt.show()
```




Figure 6: The red line is the 45 degree angle line

- Does the prediction look good? By what standard?

Predict child's heights using parent's heights

- Probably a better idea: $\hat{y} = \underbrace{\hat{\beta}_0}_{\text{intercept}} + \underbrace{\hat{\beta}_1}_{\text{slope}} x$
- The **best** fit line of this form. (what does “best” mean?)
- Why do people believe in a simple prediction formula of this kind?
 - simplicity means better interpretability potential
 - in the old days, the solution can be calculated by hand
 - OK prediction performance on some problems
- Before we continue the Python implementation of the simple linear regression, we take a detour and talk about the simple linear regression **model**
- But what is a model? More specifically, what is a simple linear regression model?

Linear regression model

- Defining the model (according to George Box)
 - All models are wrong; some models are useful...
 - Just as the ability to devise simple but evocative models is the signature of the great scientist
 - so overelaboration and overparameterization is often the mark of mediocrity (an insight when the available sample sizes were NOT huge)
- The simplest of all is a (simple) *linear* regression model; that is, the *response* or *target* y satisfies

$$y = \beta_0 + \beta_1 x + \varepsilon$$

in which

- β_0 is the *intercept* term
- β_1 is the *slope*
- β_0 and β_1 are *coefficients* or *parameters* of the linear model
- ε is a noise term, which is usually assumed independent of x and mean zero. It is often assumed to be normally distributed in theoretical analysis. Often, we assume the standard deviation of ε (σ) is unknown, which is another parameter of the simple linear regression model.

Linear regression (child heights vs. parent heights)

Suppose you have used Python to find the *best* linear fit:

$$\begin{aligned}\hat{y} &= \hat{\beta}_0 + \hat{\beta}_1 x \\ &= \hat{\beta}_0 + \hat{\beta}_1 \bar{x} + \hat{\beta}_1 \cdot (x - \bar{x}),\end{aligned}$$

where \bar{x} is the mean parent height

Why did we decompose the equation this way?

Let me first tell you that the slope estimate $\hat{\beta}_1 \in (0, 1)$ in our example

- $\hat{\beta}_0 + \hat{\beta}_1 \bar{x}$ is average child height
- $\hat{\beta}_1 \cdot (x - \bar{x})$ is *regression to mean*, so taller parents' children shrink toward average

How do $\hat{\beta}_0$ and $\hat{\beta}_1$ differ from β_0 and β_1 ?

Fitting the regression model

- **Data:** Have n pairs (x_i, y_i) , where x_i is the height of parent i and y_i is height of child i
- **Predictions:** \hat{y}_i is our model's prediction of the height of child i
- **Loss function:** First, define a way to measure *error* or *residual*
$$e_i = y_i - \hat{y}_i \qquad \text{Loss}(y, \hat{y}) = (y - \hat{y})^2$$
- Why does this loss function make sense? It is the only loss function that makes sense?

Fitting the regression model

- **Data:** Have a lot of pairs (x_i, y_i) , where x_i is the height of parent i and y_i is height of child i , and $i = 1, \dots, n$
- **Predictions:** $\hat{y}_i = \beta_0 + \beta_1 x_i$ is model prediction of child height i
- **Loss on data:** Estimate parameters β_0 and β_1 by solving least squares problem (suppose the minimizers are $\hat{\beta}_0$ and $\hat{\beta}_1$)

$$\underset{\beta_0, \beta_1}{\text{minimize}} \sum_{i=1}^n (y_i - \hat{y}_i)^2 \left(= \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i)^2 \right)$$

- $RSS = (y_1 - \hat{\beta}_0 - \hat{\beta}_1 x_1)^2 + \dots + (y_n - \hat{\beta}_0 - \hat{\beta}_1 x_n)^2$
- residual standard error (RSE)

$$RSE = \sqrt{\frac{1}{n - p - 1} RSS}$$

($p = 1$ in simple linear regression)

R^2 (a.k.a. coefficient of determination)

- To calculate R^2 , we use the formula

$$R^2 = \frac{TSS - RSS}{TSS} = 1 - \frac{RSS}{TSS}$$

where

- $RSS = \sum_{i=1}^n (y_i - \hat{y}_i)^2$ is residual sum of squares. $e_i = y_i - \hat{y}_i$ is the i th residual.
- $TSS = \sum_{i=1}^n (y_i - \bar{y})^2$ is total sum of squares
- $TSS - RSS = \sum_{i=1}^n (\hat{y}_i - \bar{y})^2$
- This R^2 definition works for both simple linear regression and multiple linear regression
- R^2 is the percent of the variation in the response explained by the regression model; a common measure for how good a linear fit is.
- $0 \leq R^2 \leq 1$ Is a bigger R^2 better?

Note: Another set of notation is also common in textbooks. Total sum of squares (TSS) is often written as SST , and residual sum of squares (RSS) is also called error sum of squares (SSE). Most confusingly, in some other textbooks, $SSR = SST - SSE$ is sum of squares due to regression. So in this alternative universe of notations, $R^2 = SSR/SST$.

R^2 (a.k.a. coefficient of determination)

- To calculate R^2 , we use the formula

$$R^2 = \frac{TSS - RSS}{TSS} = 1 - \frac{RSS}{TSS}$$

where

- $RSS = \sum_{i=1}^n (y_i - \hat{y}_i)^2$ is residual sum of squares. $e_i = y_i - \hat{y}_i$ is the i th residual.
- $TSS = \sum_{i=1}^n (y_i - \bar{y})^2$ is total sum of squares
- $TSS - RSS = \sum_{i=1}^n (\hat{y}_i - \bar{y})^2$
- This R^2 definition works for both simple linear regression and multiple linear regression
- R^2 is the percent of the variation in the response explained by the regression model; a common measure for how good a linear fit is.
- $0 \leq R^2 \leq 1$ Is a bigger R^2 better?

Note: Another set of notation is also common in textbooks. Total sum of squares (TSS) is often written as SST , and residual sum of squares (RSS) is also called error sum of squares (SSE). Most confusingly, in some other textbooks, $SSR = SST - SSE$ is sum of squares due to regression. So in this alternative universe of notations, $R^2 = SSR/SST$.

Correlation and R^2

- (Sample) correlation r between X and Y :

$$r(X, Y) = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_{i=1}^n (x_i - \bar{x})^2} \sqrt{\sum_{i=1}^n (y_i - \bar{y})^2}}$$

- Measures the **linear** dependency between two numerical variables.
- For simple linear regression, $R^2 = r^2$. This relation does not extend to multiple linear regression. Why?
- In different application fields, *good* R^2 values are vastly different.

Back to Python implementation on the height data

```
X = height.drop('child', axis =1).values
y = height['child'].values
## DataFrame.values returns the NumPy representation of the Data Frame, and the axis label will be removed.
## Equivalently, one can use .to_numpy().
print(type(X)); print(type(y)) ## try to remove the .values and see the types

<class 'numpy.ndarray'>
<class 'numpy.ndarray'>
```

Figure 7: Data preparation

```
: linear_model = LinearRegression(); linear_model.fit(X, y)
## if X were created by X = height['parent'] it is necessary to transform X
## by.reshape(-1,1) before calling the fit function, for the instructor's way,
## X is already two dimensional.
r_sq = linear_model.score(X, y); print('coefficient of determination:', r_sq)

coefficient of determination: 0.5712707984937204

: print('intercept:', linear_model.intercept_); print('slope:', linear_model.coef_)
## You can notice that .intercept_ is a scalar, while .coef_ is an array.

intercept: 25.641176413281514
slope: [0.62145545]
```

Figure 8: Fitting a simple linear regression model

Making Predictions

```
: y_pred = linear_model.predict(X) ## making prediction on the training X.

: x_new = np.arange(50, 60).reshape((-1, 1))
  ## Making prediction on some new points
  ## Here the .reshape(-1,1) function is necessary to make a 1-D array two dimensional

: y_new = linear_model.predict(x_new); print(y_new)

[56.71394872 57.33540416 57.95685961 58.57831506 59.1997705  59.82122595
 60.4426814  61.06413684 61.68559229 62.30704773]
```

Figure 9: Prediction on training and new x

But how do we find the minimizer in the loss function?

- We have use the Python `LinearRegression` as a blackbox to find the $\hat{\beta}_0$ and $\hat{\beta}_1$ that minimize the RSS.
- People in the precomputer age (or in exam settings) use an exact formula

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2} \text{ and } \hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x}$$

- remark: every least squares regression line passes (\bar{x}, \bar{y})
- Modern packages use optimization techniques. Not covered in DSO 530.