

# DSO530 Statistical Learning Methods

## Lecture 7b : Bagging, Random Forests and Boosting

Dr. Xin Tong

Department of Data Sciences and Operations

Marshall School of Business

University of Southern California

## Review: Advantages and disadvantages of trees

- Trees are very easy to explain to people. In fact, they are even easier to explain than linear regression!
- Why? No mathematical formula needed in the communication
- Some people believe that decision trees more closely mirror human decision-making than do the regression and classification approaches seen in previous lectures.
- Trees can be displayed graphically, and are easily interpreted even by a non-expert (especially if they are small).
- **Unfortunately**, trees generally do not have the great predictive accuracy.
- This disadvantage motivates the ensemble methods such as *bagging*, *random forests* and *boosting*.

# Bagging

- The decision trees typically suffer from high variance.
- *Bootstrap aggregation*, or **bagging**, is a general-purpose procedure for reducing the variance of a statistical learning method; we introduce it here because it is particularly useful and frequently used in the context of decision trees.
- In this approach we generate  $B$  different *bootstrapped* training data sets. We then train our method on the  $b$ th bootstrapped training set in order to get  $\hat{f}^{*b}(x)$  (not pruned), and average all the predictions:

$$\hat{f}_{bag}(x) = \frac{1}{B} \sum_{b=1}^B \hat{f}^{*b}(x).$$


- Averaging these  $B$  trees reduces the variance. (Why?)
- $B$  is not a critical parameter with bagging; a very large value of  $B$  will not lead to overfitting.
- In practice, use  $B$  sufficiently large so that the error has settled down.
- For a given test observation in `classification`, we can record the class predicted by each of the  $B$  trees, and take a vote: the overall prediction is the most commonly occurring class among the  $B$  predictions.

# Out-of-bag Error Estimation

- For a bagged model, we can estimate the test error without doing CV or a validation set approach
- Each bagged tree makes use of a part of the original observations
- The remaining observations not used to fit a given bagged tree are referred to as the *out-of-bag* (OOB) observations
- We can predict the response for the  $i$ th observation using each of the trees in which that observation was OOB
- This will yield a little more than  $B/3$  predictions for the  $i$ th observation (what is  $(1 - 1/n)^n$  as  $n$  goes to infinity? )
- To obtain a single prediction for the  $i$ th observation, we can average these predicted responses (if regression is the goal) or can take a majority vote (if classification is the goal). This leads to a single OOB prediction for the  $i$ th observation
- An OOB prediction can be obtained in this way for each of the  $n$  observations, from which the overall OOB MSE (for a regression problem) or classification error (for a classification problem) can be computed

# Variable Importance Measures

- Although the collection of bagged trees is much more difficult to interpret than a single tree, one can obtain an overall summary of the importance of each predictor using
  - the RSS (for bagging regression trees)
  - the Gini index (for bagging classification trees)
- In the case of bagging regression trees, we can record the total amount that the RSS is decreased due to splits over a given predictor, averaged over all  $B$  trees. A large value indicates an important predictor
- For bagging classification trees, we can add up the total amount that the Gini index is decreased by splits over a given predictor, averaged over all  $B$  trees.



**FIGURE 8.9.** *A variable importance plot for the **Heart** data. Variable importance is computed using the mean decrease in Gini index, and expressed relative to the maximum.*

Figure 1

# Random Forest(s)

- Bagging constructs trees that are too “similar” (why?), so it probably does not reduce the variance as much as we wish to.
- **Random forests** provide an improvement over bagged trees by a small tweak that *decorrelates* the trees.
- As in bagging, we build a number of decision trees on bootstrapped training samples.
- But when building these decision trees, each time a split in a tree is considered, *a random sample of  $m$  predictors* is chosen as split candidates from the full set of  $p$  predictors. The split is allowed to use *only one* of those  $m$  predictors.
- So bagging is a special case of random forest when  $m = p$
- If one does not want to spend extra efforts on  $m$ , one might use  $m = \sqrt{p}$  as a canonical choice for classification and  $m = p/3$  as a canonical choice for regression.
- As with bagging, random forests will not overfit if we increase  $B$ , so in practice people use  $B$  sufficiently large for the error rate to have settled down.
- Random forest is a really good off-the-shelf algorithm




Figure 2



# Python implementation

- `RandomForestClassifier` and `RandomForestRegressor` in `sklearn` implement random forests in Python for classification and regression problems, respectively
- Our tutorial covers `RandomForestClassifier`
- Parameters: `n_estimators` (default 100) is the number of trees in the forest; `max_features` (default `sqrt(n_features)`) is the number of features to consider when looking for the best split.
- You can learn pick up `RandomForestRegressor` from <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestRegressor.html>
- `RandomForestRegressor` in `sklearn` has a default setting of `max_features=n_features`. This is at odds with the recommendation by ISLR

# A definition and some questions

- If you need to communicate a one sentence ad-hoc definition of random forests:
  - *Random forests are bagged decision tree models that split on a random subset of features on each split.*
- Q: In random forest algorithms, we restrict our attention to randomly selected  $m$  out of  $p$  features in each split. Now we change this procedure to restriction to the first  $m$  features (i.e.,  $X_1, \dots, X_m$ ) in every split. Do you expect the new procedure to work well? And why?
- Q: If a decision tree partitions the feature space into regions  $R_1, \dots, R_J$ , can any of these regions be a ball?
- Q: Is random forest always a better algorithm compared to decision trees?
- Q: What are the sources of randomness that a random forest model has? Hint: 3.
- Q: (open question) Can we extend the random forest idea to other base algorithms rather than trees?

# A definition and some questions

- If you need to communicate a one sentence ad-hoc definition of random forests:
  - *Random forests are bagged decision tree models that split on a random subset of features on each split.*
- Q: In random forest algorithms, we restrict our attention to randomly selected  $m$  out of  $p$  features in each split. Now we change this procedure to restriction to the first  $m$  features (i.e.,  $X_1, \dots, X_m$ ) in every split. Do you expect the new procedure to work well? And why?
- Q: If a decision tree partitions the feature space into regions  $R_1, \dots, R_J$ , can any of these regions be a ball?
- Q: Is random forest always a better algorithm compared to decision trees?
- Q: What are the sources of randomness that a random forest model has? Hint: 3.
- Q: (open question) Can we extend the random forest idea to other base algorithms rather than trees?

# Boosting

- Like bagging, **boosting** is a general approach that can be applied to many statistical learning methods for regression or classification.
- Boosting is an ensemble technique where new models are added to correct the errors made by existing models.
- A differentiating characteristic Random forest: parallel vs. boosting: sequential

# An example from the ISLR book

---

**Algorithm 8.2** *Boosting for Regression Trees*

---

1. Set  $\hat{f}(x) = 0$  and  $r_i = y_i$  for all  $i$  in the training set.
2. For  $b = 1, 2, \dots, B$ , repeat:
  - (a) Fit a tree  $\hat{f}^b$  with  $d$  splits ( $d + 1$  terminal nodes) to the training data  $(X, r)$ .
  - (b) Update  $\hat{f}$  by adding in a shrunk version of the new tree:

$$\hat{f}(x) \leftarrow \hat{f}(x) + \lambda \hat{f}^b(x). \quad (8.10)$$

- (c) Update the residuals,

$$r_i \leftarrow r_i - \lambda \hat{f}^b(x_i). \quad (8.11)$$

3. Output the boosted model,

$$\hat{f}(x) = \sum_{b=1}^B \lambda \hat{f}^b(x). \quad (8.12)$$

---

Figure 3: Boosting Tree

# In the boosting tree algorithm

- We use cross-validation to select  $B$ .
- The shrinkage parameter  $\lambda$ , a small positive number. This controls the rate at which boosting learns. Typical values are 0.01 or 0.001.
- The number  $d$  of splits in each tree, which controls the complexity of the boosted ensemble. Often  $d = 1$  works well, in which case each tree is a *stump*, consisting of a single split. More generally  $d$  is the *interaction depth*.
- Boosting is a *slow learner*.

## In practice

- XGBoost is the to-go implementation of boosting algorithms for its execution speed and model performance.
- It is more complicated than what we described on the previous slide. For example, subsampling and shrinkage ideas are adopted.
- `xgboost` available in Python.