DSO530 Statistical Learning Methods

Lecture 9 : Unsupervised Learning part I

Dr. Xin Tong
Department of Data Sciences and Operations
Marshall School of Business
University of Southern California

Introduction

- Unsupervised Learning Topics include
  - *principal components analysis* (I): for data visualization or data pre-processing before supervised techniques are applied
  - *clustering* (II): or discovering unknown subgroups in data
- Recall the difference between supervised learning and unsupervised learning (Labeling is usually costly)
- Name a few unsupervised learning examples?
- It is not so obvious to assess unsupervised learning
- Unsupervised learning is developing fast these days. You should not think that the methods covered in ISLR are all of the state-of-the-art

Introduction

- Unsupervised Learning Topics include
  - *principal components analysis* (I): for data visualization or data pre-processing before supervised techniques are applied
  - *clustering* (II): or discovering unknown subgroups in data
- Recall the difference between supervised learning and unsupervised learning (Labeling is usually costly)
- Name a few unsupervised learning examples?
- It is not so obvious to assess unsupervised learning
- Unsupervised learning is developing fast these days. You should not think that the methods covered in `ISLR` are all of the state-of-the-art

Principal Components Analysis

- *Principal components* were discussed in the context of principal components regression
- The principal component directions are directions in feature space along which the original data are highly variable
- Principal component analysis (PCA) refers to the process by which principal components are computed, and the subsequent use of these components in understanding the data
- If we were to do scatterplot for every pair of variables when $p = 10$, how many scatterplot do we need?
- Clearly, a better method is required to visualize the n observations when p is large

- First principal component of a set of features $X_1, X_2, \cdots, X_p$ is the *normalized* linear combination of the features

$$Z_1 = \phi_{11} X_1 + \phi_{21} X_2 + \cdots + \phi_{p1} X_p \,, \text{ where } \sum_{j=1}^{p} \phi_{j1}^2 = 1 \,.$$

- We refer to $\phi_{11}, \cdots, \phi_{p1}$ as the loadings of the first principal component.
- Given a $n \times p$ data set X, how do we compute the first principal component?
- Assume that each of the variables in X has been centered to have mean zero
- We then look for the linear combination of the sample feature values of the form

$$z_{i1} = \phi_{11} x_{i1} + \cdots + \phi_{21} x_{i2} + \cdots + \phi_{p1} x_{ip}$$

that has largest sample variance, subject to the constraint that $\sum_{j=1}^{p} \phi_{j1}^2 = 1$

- In other words, the first principal component loading vector solves the optimization problem

$$\max_{\phi_{11},\cdots,\phi_{p1}} \left\{ \frac{1}{n} \sum_{i=1}^{n} z_{i1}^2 \right\}, \text{ subject to } \sum_{j=1}^{p} \phi_{j1}^2 = 1 .$$

- This can be solved by an eigen decomposition, but this is beyond the scope of the book
- We refer to $z_{11}, \cdots, z_{n1}$ as the scores of the first principal component
- The loading vector $\phi_1 = (\phi_{11}, \cdots, \phi_{p1})^T$ defines a direction in feature space along which the data vary the most
- If we project n data points $x_1, \cdots, x_n$ on to ϕ_1, the projected values are $z_{11}, \cdots, z_{n1}$
- The second principal component is the linear combination of $X_1, \cdots, X_p$ that has *maximal* variance out of all linear combinations that are *uncorrelated* with the first component Z_1.

- The second principal components scores $z_{12}, z_{22}, ..., z_{n2}$ take the form

$$z_{i2} = \phi_{12}x_{i1} + \phi_{22}x_{i2} + \cdots + \phi_{p2}x_{ip}$$

- "Z_2 and Z_1 are uncorrelated" is equivalent to "ϕ_1 is perpendicular to ϕ_2"
- [Optional] The principal component directions $\phi_1, \phi_2, \phi_3, \cdots$ are the ordered sequence of *eigenvectors* of the matrix X^TX, and the variances of the components are the *eigenvalues*. An equivalent formulation is through the so-called *singular value decomposition* (SVD) of X
- The maximum number of PCs is $\min(n-1, p)$ (think about 2 points on a plane)

	PC1	PC2
Murder	0.5358995	−0.4181809
Assault	0.5831836	−0.1879856
UrbanPop	0.2781909	0.8728062
Rape	0.5434321	0.1673186

TABLE 10.1. *The principal component loading vectors, ϕ_1 and ϕ_2, for the* `USArrests` *data. These are also displayed in Figure 10.1.*

Figure 1

- A *biplot* (optional): a figure that represents both the principal component scores and the loading vectors

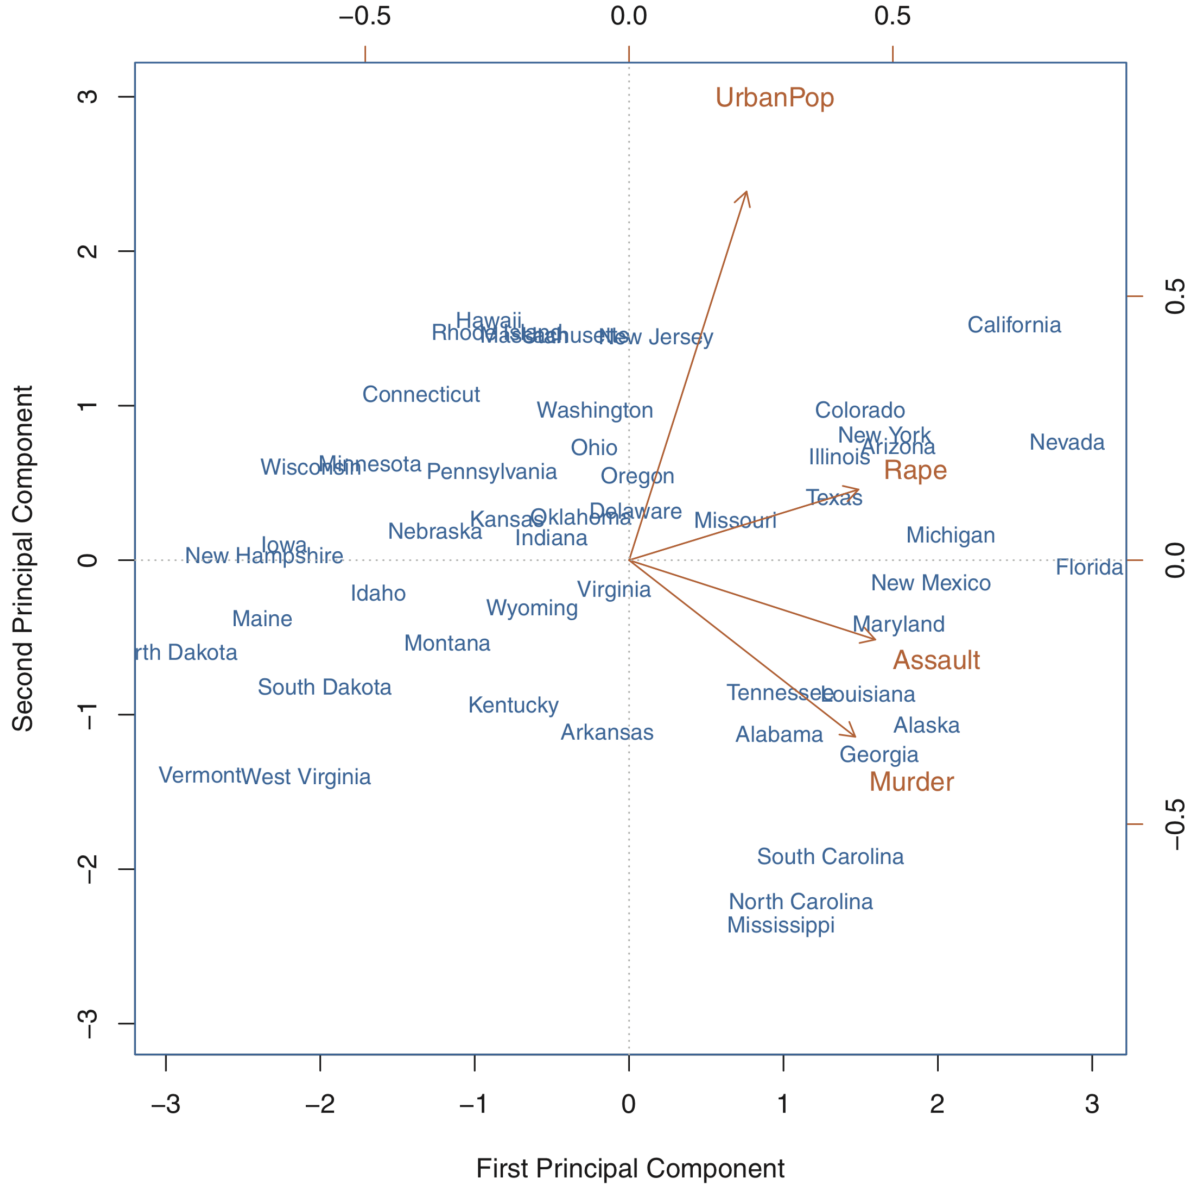

Another Interpretation of Principal Components (Optional)

- Principal components provide low-dimensional linear surfaces that are *closest* to the observations
- The first principal component loading vector represent the line in p-dimensional space that is closest to the n observations
- The first two principal components of a data set span the plane that is closest to the n observations
- $x_{ij} \approx \sum_{m=1}^{M} z_{im}\phi_{jm}$
- When $M = \min(n-1, p)$, $x_{ij} = \sum_{m=1}^{M} z_{im}\phi_{jm}$

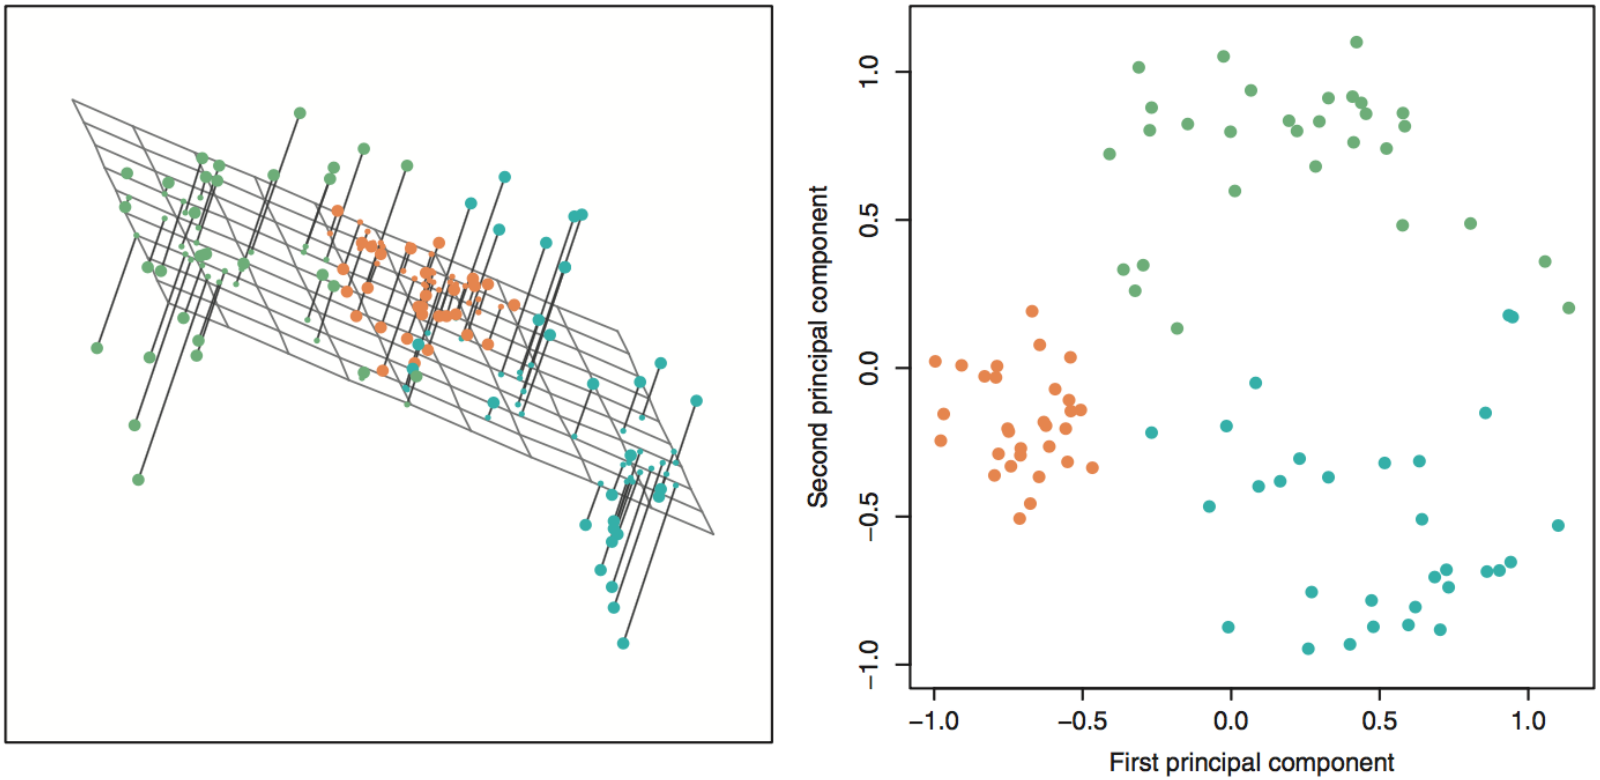

FIGURE 10.2. *Ninety observations simulated in three dimensions.* Left: *the first two principal component directions span the plane that best fits the data. It minimizes the sum of squared distances from each point to the plane.* Right: *the first two principal component score vectors give the coordinates of the projection of the 90 observations onto the plane. The variance in the plane is maximized.*

Figure 2

- Usually we standardize variables (mean 0 and SD 1) before doing PCA

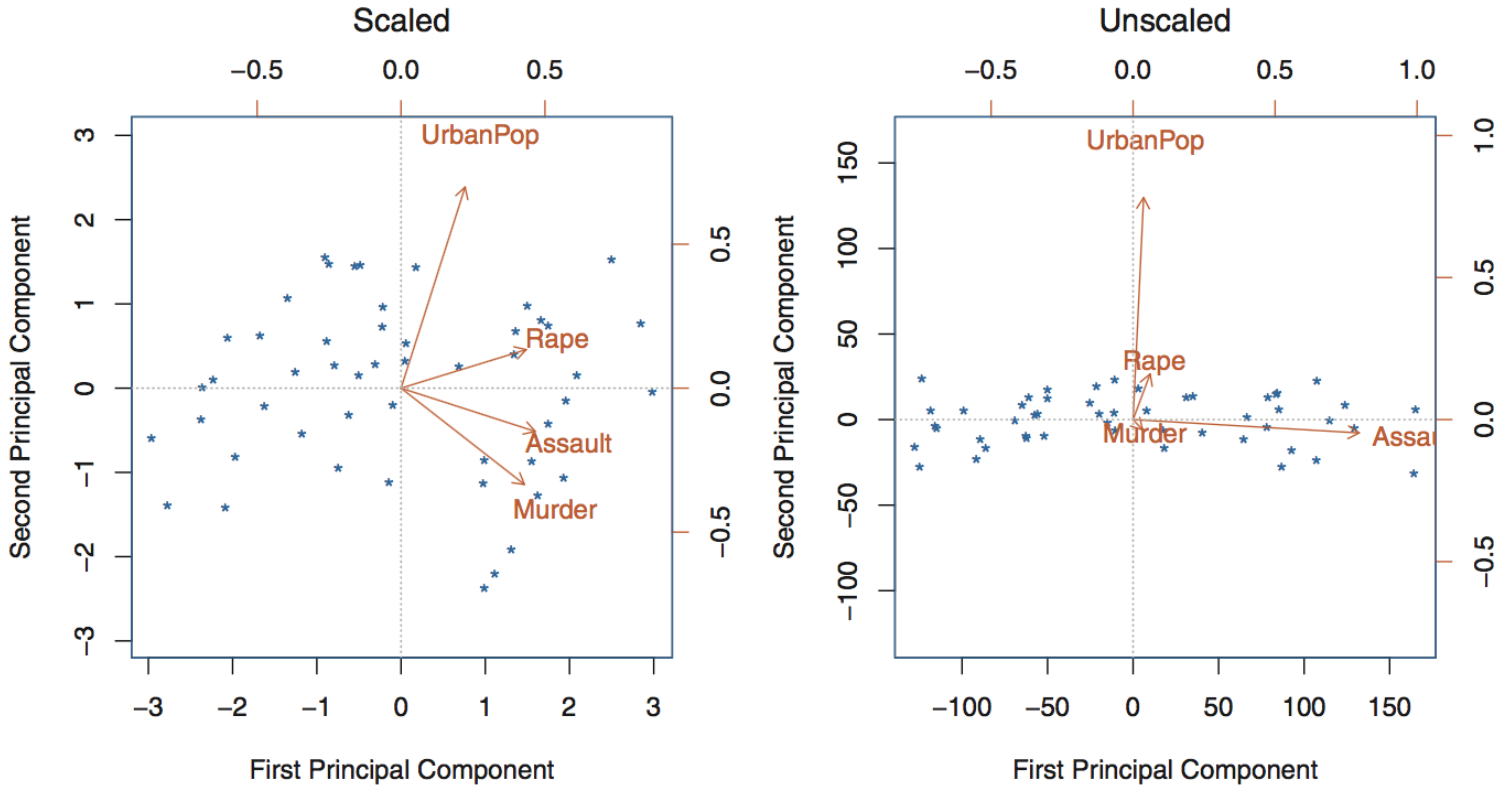

FIGURE 10.3. *Two principal component biplots for the* `USArrests` *data.* Left: *the same as Figure 10.1, with the variables scaled to have unit standard deviations.* Right: *principal components using unscaled data.* `Assault` *has by far the largest loading on the first principal component because it has the highest variance among the four variables. In general, scaling the variables to have standard deviation one is recommended.*

- In certain settings, variables may be measured in the same units. Then, we might choose not to scale the variables before PCA

Proportion of Variance Explained

- We are interested in knowing the *proportion of variance explained* (PVE) by each principal component
- The total variance present in a data set (assuming that the variables have been centered to have mean zero) is defined as

$$\sum_{j=1}^{p} \frac{1}{n} \sum_{i=1}^{n} x_{ij}^2$$

- The variance explained by the mth principal component

$$\frac{1}{n} \sum_{i=1}^{n} z_{im}^2 = \frac{1}{n} \sum_{i=1}^{n} \left(\sum_{j=1}^{p} \phi_{jm} x_{ij} \right)^2$$

- The PVE of the mth principal component is given by

$$\frac{\sum_{i=1}^{n} \left(\sum_{j=1}^{p} \phi_{jm} x_{ij} \right)^2}{\sum_{j=1}^{p} \sum_{i=1}^{n} x_{ij}^2}$$

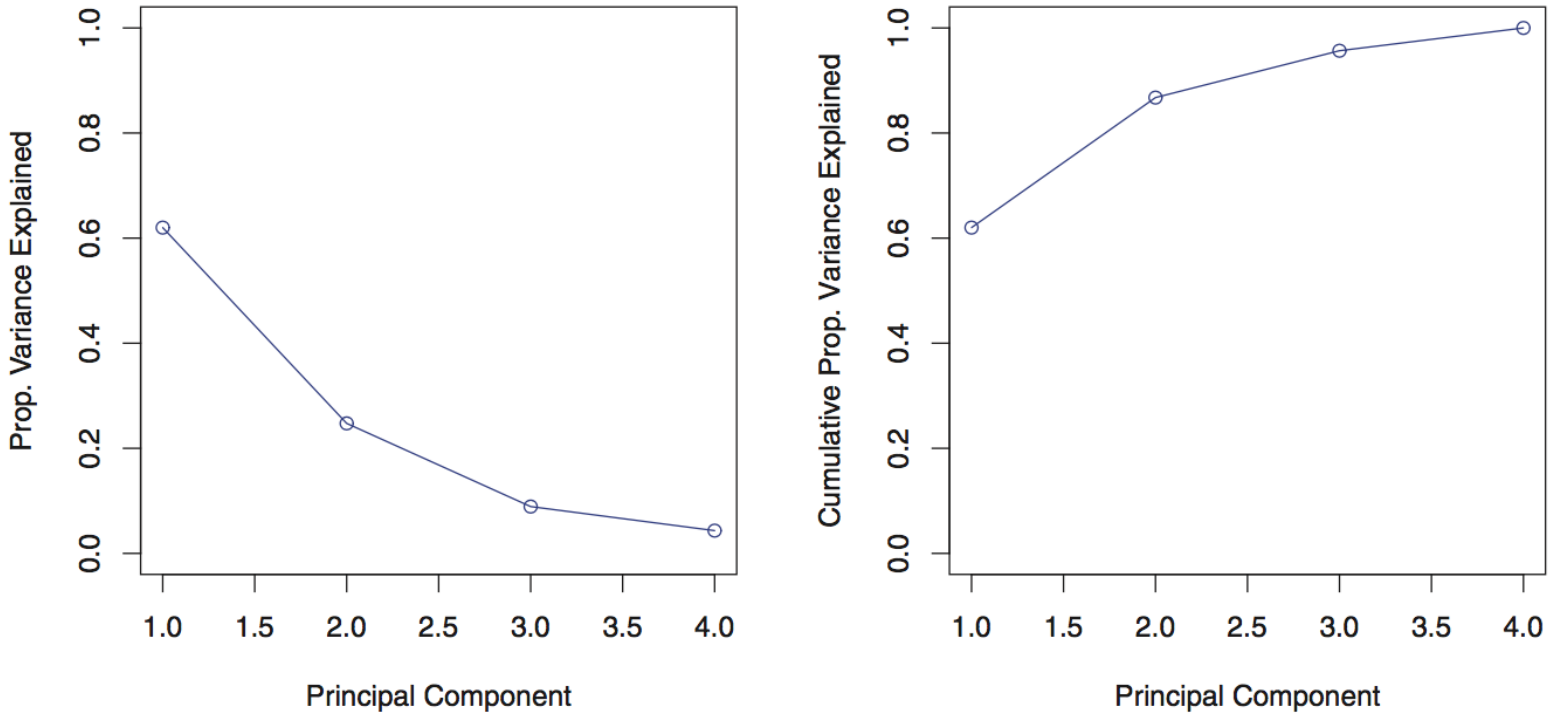

FIGURE 10.4. Left: *a scree plot depicting the proportion of variance explained by each of the four principal components in the* `USArrests` *data.* Right: *the cumulative proportion of variance explained by the four principal components in the* `USArrests` *data.*

Figure 3

- The question of how many principal components are enough is inherently ill-defined, and will depend on the specific area of application and the specific data set