

# Week 1 Lab: Introduction to SQL

## Parch and Posey Database

### Contents

Creating the Database . . . . .	1
Quering the Database . . . . .	3
SELECT & FROM . . . . .	3
LIMIT . . . . .	5
ORDER BY . . . . .	6
WHERE . . . . .	7
Derived Columns . . . . .	9
Logical Operators . . . . .	9

Parch & Posey is a fictional company that sells paper:

- There are 50 sales reps spread across the United States in 4 regions.
- There are 3 types of paper. Regular, Poster and Glossy.
- The clients are primarily large Fortune's 100 companies.

The Parch & Posey database is provided by Mode Analytics (<https://mode.com/>).

Questions answered using Parch & Posey data are meant to simulate real word problems. Using SQL, we will help Parch & Posey answer questions like: Which of their product lines is worst performing? Which of their market channels they should make a great investment in.

In the Parch & Posey database there are five tables (essentially 5 spreadsheets):

- web\_events
- accounts
- orders
- sales\_reps
- region

Figure 1 shows the ERD (entity relationship diagram) for Parch and Posey.

Most of the variables in each table are self-explanatory, but some are not. Here is a description of them:

- channel: marketing channel (twitter, adwords, organic, banner, facebook, direct)
- lat: latitude of the company location
- long: longitude of the company location
- primary\_poc: primary point of contact
- sales\_rep\_id: sales representative id
- standard\_qty: quantity of standard papers ordered
- poster\_qty: quantity of poster papers ordered
- glossy\_qty: quantity of glossy papers ordered
- total: total quantity of papers ordered (standard + poster + glossy)
- standard\_amt\_usd: dollar amount paid for the standard papers
- poster\_amt\_usd: dollar amount paid for the poster papers
- gloss\_amt\_usd: dollar amount paid for the gloss papers
- total\_amt\_usd: dollar amount paid for all paper orders

### Creating the Database

- Open pgAdmin 4

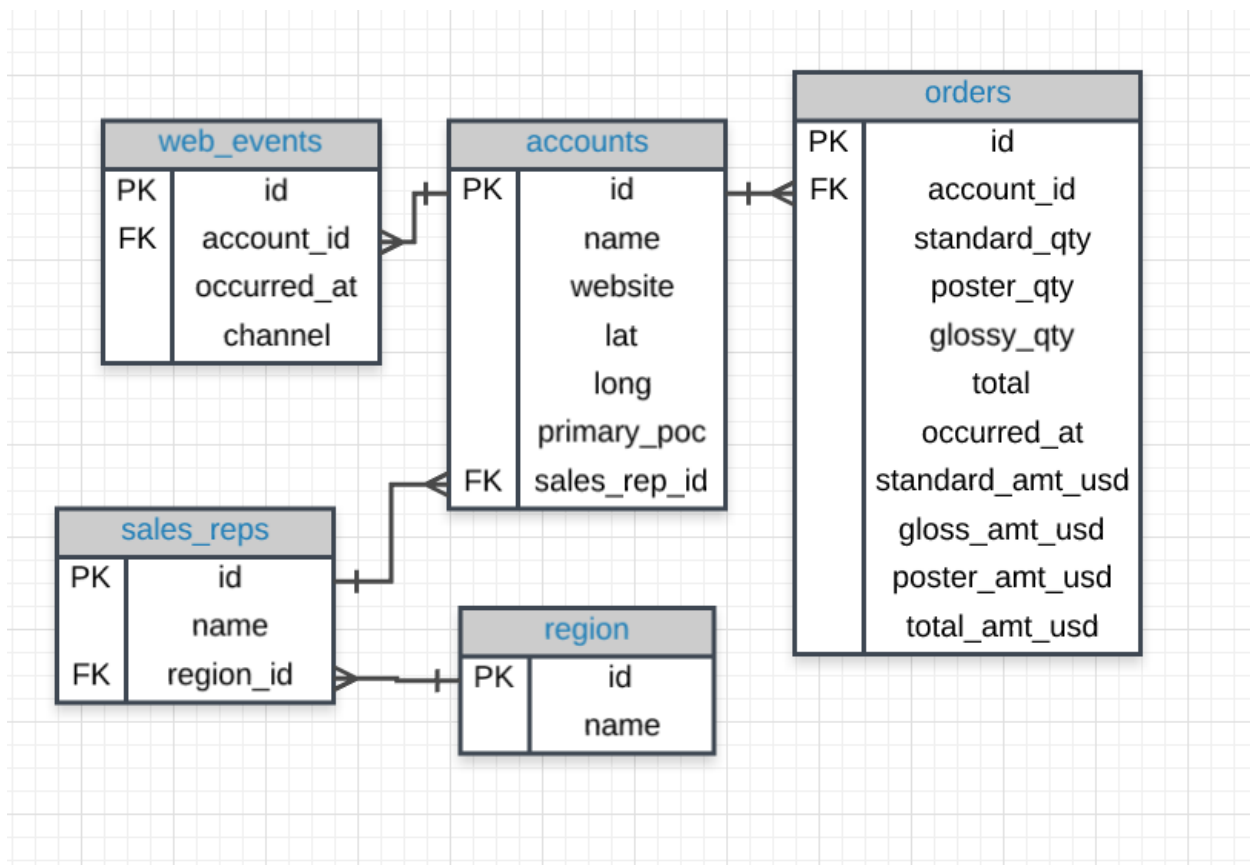


Figure 1: Parch and Posey ERD

- Create a new database, and give it a name. In our case, let's call it "Parch".
- Right click on the Parch database, and choose "Query Tool..."
- Load the database file "parch.sql" and run all the commands (We will discuss these commands later during the semester, but for now, just think of this step as a database creation and connection)

## Querying the Database

To query the database, we need a SQL IDE (Interface Development Environment). There are many IDE's that are compatible with PostgreSQL, and feel free to use any!

In this section I will introduce two IDEs. The first is pgAdmin, and the second is Rstudio.

If you decide to stick to pgAdmin, then all you have to do (after creating the database) is to \*\* Right click on the Parch database, and choose "Query Tool..." to create an empty SQL file. \*\*

If you decide to use the Rstudio IDE, then follow these steps:

- Open pgAdmin 4 and connect to the database of your choice. In this lab, we will connect with the Parch database.
- Once connected, open Rstudio, load the library RPostgreSQL using the following command (Initially, you will need to install the R package).

```
# use the following command to install the RPostgreSQL package
# install.packages("RPostgreSQL")
```

```
# load the RPostgreSQL package using the Console
library(RPostgreSQL)
```

Once the package is loaded, then open a new SQL script file from the File menu in Rstudio. This will open an empty file with a line at the top. Change this file to:

```
# Change the dbname to "Parch" in our case (name of the database)
# Change the password to your PostgreSQL password
```

```
-- !preview conn=DBI::dbConnect(DBI::dbDriver("PostgreSQL"), dbname = "DATABASENAME", host="localhost", p
```

You're all set!

## SELECT & FROM

1. Take a quick look at all the 5 tables in the database.

```
SELECT *
FROM sales_reps
limit 5;
```

id	name	region_id
321500	Samuel Racine	1
321510	Eugena Esser	1
321520	Michel Averette	1
321530	Renetta Carew	1
321540	Cara Clarke	1

```
SELECT *
```

```
FROM orders
limit 5;
```

id	account_id	occurred_at	standard_qty	gross_qty	poster_qty	total	standard_amt_usd	gross_amt_usd	poster_amt_usd	total_amt_usd
1	1001	2015-10-06 17:31:14	123	22	24	169	613.77	164.78	194.88	973.43
2	1001	2015-11-05 03:34:33	190	41	57	288	948.10	307.09	462.84	1718.03
3	1001	2015-12-04 04:21:55	85	47	0	132	424.15	352.03	0.00	776.18
4	1001	2016-01-02 01:18:24	144	32	0	176	718.56	239.68	0.00	958.24
5	1001	2016-02-01 19:27:27	108	29	28	165	538.92	217.21	227.36	983.49

```
SELECT *
FROM accounts
limit 5;
```

id	name	website	lat	long	primary_poc	sales_rep_id
1001	Walmart	www.walmart.com	40.23850	-75.10330	Tamara Tuma	321500
1011	Exxon Mobil	www.exxonmobil.com	41.16916	-73.84937	Sung Shields	321510
1021	Apple	www.apple.com	42.29049	-76.08401	Jodee Lupo	321520
1031	Berkshire Hathaway	www.berkshirehathaway.com	40.94902	-75.76390	Serafina Banda	321530
1041	McKesson	www.mckesson.com	42.21709	-75.28500	Angeles Crusoe	321540

```
SELECT *
FROM region
limit 5;
```

id	name
1	Northeast
2	Midwest
3	Southeast
4	West

```
SELECT *
FROM web_events
limit 5;
```

id	account_id	occurred_at	channel
1	1001	2015-10-06 17:13:58	direct
2	1001	2015-11-05 03:08:26	direct
3	1001	2015-12-04 03:57:24	direct
4	1001	2016-01-02 00:55:03	direct
5	1001	2016-02-01 19:02:33	direct

2. Generate a list of all the names of the sales representatives and their IDs that Parch and Posey have in their database.

To answer this question, we will use a SELECT statement. The SELECT statement is used to query the database and retrieve selected data that match the criteria that you specify. In this case, we are looking to retrieve only the sales reps **name** and **id**, so we just choose these two columns in the SELECT statement.

The FROM statement is where you tell the query what table you are querying from. In this case, we are looking to get the information from the **sales\_reps** table in the database.

```
SELECT id, name
```

```
FROM sales_reps
limit 5;
```

The following output only shows the first few records on the output.

id	name
321500	Samuel Racine
321510	Eugena Esser
321520	Michel Averette
321530	Renetta Carew
321540	Cara Clarke

3. **YOUR TURN** Write your own query to select only the `id`, `account_id`, and `occurred_at` columns for all orders in the `orders` table.

```
SELECT id, account_id, occurred_at
FROM orders
limit 5;
```

The following output only shows the first few records on the output.

id	account_id	occurred_at
1	1001	2015-10-06 17:31:14
2	1001	2015-11-05 03:34:33
3	1001	2015-12-04 04:21:55
4	1001	2016-01-02 01:18:24
5	1001	2016-02-01 19:27:27

Sometimes, and for the purpose of viewing the results, we might decide to limit our output to a certain number of rows/tuple/observations. For this purpose, we can use the `LIMIT` command. Please note that the `LIMIT` command is always the very last part of a query.

## LIMIT

4. Showing just the first 10 observations of the `sales_reps` table with all of the columns.

```
SELECT *
FROM sales_reps
LIMIT 10;
```

id	name	region_id
321500	Samuel Racine	1
321510	Eugena Esser	1
321520	Michel Averette	1
321530	Renetta Carew	1
321540	Cara Clarke	1
321550	Lavera Oles	1
321560	Elba Felder	1
321570	Shawanda Selke	1
321580	Sibyl Lauria	1
321590	Necole Victory	1

## ORDER BY

The ORDER BY statement allows us to order our table by any row. If you are familiar with Excel, this is similar to the sorting you can do with filters.

The ORDER BY statement is always after the SELECT and FROM statements, but it is before the LIMIT statement. As you learn additional commands, the order of these statements will matter more. If we are using the LIMIT statement, it will always appear last.

5. Write a SQL query to look up the most 10 recent orders.

The Order By clause will help you accomplish this by allowing you to sort the orders by date. The orders table is sorted by Account ID by default. Let's add an Order By clause to reorder the results based on the date the order was placed, which you can see in the occurred\_at column. Notice that the order by clause goes between the From and Limit clauses. You have to write the clauses in this order, or the query will not run. By default, order by goes from a to z, lowest to highest, or earliest to latest, if working with dates. If you want to order the other way, you can add DESC, short for descending, to the end of the Order By clause.

```
SELECT *
FROM orders
ORDER BY occurred_at DESC
LIMIT 10;
```

id	account_id	occurred_at	standard_qty	gloss_qty	poster_qty	total	standard_amt_usd	gloss_amt_usd	poster_amt_usd	total_amt_usd
6451	3841	2017-01-02 00:02:40	42	506	302	850	209.58	3789.94	2452.24	6451.76
3546	3841	2017-01-01 23:50:16	291	36	26	353	1452.09	269.64	211.12	1932.85
6454	3861	2017-01-01 22:29:50	38	167	51	256	189.62	1250.83	414.12	1854.57
3554	3861	2017-01-01 22:17:26	497	0	23	520	2480.03	0.00	186.76	2666.79
6556	4051	2017-01-01 21:04:25	0	65	50	115	0.00	486.85	406.00	892.85
3745	4051	2017-01-01 20:52:23	495	15	0	510	2470.05	112.35	0.00	2582.40
1092	1761	2017-01-01 17:34:10	62	28	124	214	309.38	209.72	1006.88	1525.98
1364	1961	2017-01-01 16:40:57	102	39	29	170	508.98	292.11	235.48	1036.57
3159	3431	2017-01-01 14:05:39	302	29	18	349	1506.98	217.21	146.16	1870.35
6223	3431	2017-01-01 13:57:21	51	444	135	630	254.49	3325.56	1096.20	4676.25

6. **YOUR TURN** Write a query to return the 10 earliest orders in the orders table. Include the id, occurred\_at, and total\_amt\_usd.

```
SELECT id, occurred_at, total_amt_usd
FROM orders
ORDER BY occurred_at
LIMIT 10;
```

id	occurred_at	total_amt_usd
5786	2013-12-04 04:22:44	627.48
2415	2013-12-04 04:45:54	2646.77
4108	2013-12-04 04:53:25	2709.62
4489	2013-12-05 20:29:16	277.13
287	2013-12-05 20:33:56	3001.85
1946	2013-12-06 02:13:20	2802.90
6197	2013-12-06 12:55:22	7009.18
3122	2013-12-06 12:57:41	1992.13
6078	2013-12-06 13:14:47	6680.06
2932	2013-12-06 13:17:25	2075.94

7. **YOUR TURN** Write a query to return the date of the 5 most expensive orders of papers. Make sure to include the order id, and the total dollar amount.

```
SELECT id, occurred_at, total_amt_usd
```

```
FROM orders
ORDER BY total_amt_usd DESC
LIMIT 5;
```

id	occurred_at	total_amt_usd
4016	2016-12-26 08:53:24	232207.07
3892	2016-06-24 13:32:55	112875.18
3963	2015-03-30 00:05:30	107533.55
5791	2014-10-24 12:06:22	95005.82
3778	2016-07-17 14:50:43	93547.84

8. Write a query that returns all accounts sorted by account id. For each account, list the orders sorted by total dollar amount in descending order.

You can also order by multiple columns. This is particularly useful if your data falls into categories and you'd like to organize rows by date, for example, but keep all of the results within a given category together. The statement sorts according to columns listed from left first and those listed on the right after that. We still have the ability to flip the way we order using DESC.

```
SELECT account_id, total_amt_usd
FROM orders
ORDER BY account_id, total_amt_usd DESC
LIMIT 5;
```

account_id	total_amt_usd
1001	9426.71
1001	9230.67
1001	9134.31
1001	8963.91
1001	8863.24

9. **YOUR TURN** Write a query that returns the top 5 rows with the highest total number of orders. If two rows have the same number of total orders, then use the total dollar amount (highest to lowest) to break the tie. (Return the ID, the total number of orders, and the total dollar amount)

```
SELECT id, total, total_amt_usd
FROM orders
ORDER BY total DESC, total_amt_usd DESC
LIMIT 5;
```

id	total	total_amt_usd
4016	28799	232207.07
3892	22610	112875.18
4562	16428	84099.62
3963	14395	107533.55
731	12598	92991.05

## WHERE

Imagine yourself as an account manager at Parch and Posey. You're about to head out to visit one of your most important customers and you want to show up prepared, which means making sure that you are up to speed on all of their recent purchases.

You can use a WHERE clause to generate a list of all purchases made by that specific customer. The WHERE clause allows you to filter a set of results based on specific criteria as you would with Excel's filter capability.

The WHERE clause goes after FROM but before ORDER BY or LIMIT. In addition, the following are the comparison operators with their SQL syntax:

Operator	SQL Syntax
Equal	=
Greater Than	>
Less Than	<
Great Than or Equal	>=
Less Than or Equal	<=
Not Equal	<> or !=
String Comparison Test	LIKE

10. Write a query to show only orders from our customer with an account ID 4251.

```
SELECT *
FROM orders
WHERE account_id = 4251;
```

id	account_id	occurred_at	standard_qty	gloss_qty	poster_qty	total	standard_amt_usd	gloss_amt_usd	poster_amt_usd	total_amt_usd
4009	4251	2016-06-05 01:36:42	626	15	0	641	3123.74	112.35	0.00	3236.09
4010	4251	2016-07-04 12:34:49	498	6	2	506	2485.02	44.94	16.24	2546.20
4011	4251	2016-08-02 00:53:28	679	36	5	720	3388.21	269.64	40.60	3698.45
4012	4251	2016-09-01 02:32:51	503	13	32	548	2509.97	97.37	259.84	2867.18
4013	4251	2016-09-30 13:31:53	503	39	0	542	2509.97	292.11	0.00	2802.08
4014	4251	2016-10-29 12:04:06	483	12	0	495	2410.17	89.88	0.00	2500.05
4015	4251	2016-11-27 15:17:06	520	21	7	548	2594.80	157.29	56.84	2808.93
4016	4251	2016-12-26 08:53:24	521	16	28262	28799	2599.79	119.84	229487.44	232207.07
6719	4251	2016-06-05 01:16:37	0	78	0	78	0.00	584.22	0.00	584.22
6720	4251	2016-08-02 01:13:08	9	0	19	28	44.91	0.00	154.28	199.19
6721	4251	2016-09-01 02:39:55	3	71	50	124	14.97	531.79	406.00	952.76
6722	4251	2016-11-27 15:16:07	19	56	0	75	94.81	419.44	0.00	514.25
6723	4251	2016-12-26 08:39:56	0	31	21	52	0.00	232.19	170.52	402.71

11. **YOUR TURN** Write a query to pull the first 5 rows and all columns from the orders table that have spent a total amount less than or equal to \$500.

```
Select *
FROM orders
WHERE total_amt_usd <= 500
LIMIT 5;
```

id	account_id	occurred_at	standard_qty	gloss_qty	poster_qty	total	standard_amt_usd	gloss_amt_usd	poster_amt_usd	total_amt_usd
67	1091	2015-04-07 13:29:20	95	0	0	95	474.05	0.00	0.00	474.05
96	1101	2016-03-15 11:36:03	14	8	16	38	69.86	59.92	129.92	259.70
119	1131	2016-06-12 12:29:45	0	30	23	53	0.00	224.70	186.76	411.46
124	1131	2016-11-07 05:10:56	0	0	0	0	0.00	0.00	0.00	0.00
254	1251	2014-11-01 02:15:24	0	0	17	17	0.00	0.00	138.04	138.04

12. For the account Exxon Mobil, return the the company name, website, and the primary point of contact (primary\_poc).

```
select name, website, primary_poc
from accounts
where name = 'Exxon Mobil';
```

name	website	primary_poc
Exxon Mobil	www.exxonmobil.com	Sung Shields



## Derived Columns

Creating a new column that is a combination of existing columns is known as a derived column. Derived columns can include simple arithmetic or any number of advanced calculations.

13. For each order, return the quantity of non-standard papers (poster and gloss). Include the `account_id` and date.

```
select account_id, occurred_at, gloss_qty, poster_qty, gloss_qty + poster_qty
from orders
limit 5;
```

account_id	occurred_at	gloss_qty	poster_qty	?column?
1001	2015-10-06 17:31:14	22	24	46
1001	2015-11-05 03:34:33	41	57	98
1001	2015-12-04 04:21:55	47	0	47
1001	2016-01-02 01:18:24	32	0	32
1001	2016-02-01 19:27:27	29	28	57

We can give the new derived column an alias and we can do this by adding AS to the end of the line that produces the derived column and then giving it a name.

```
select account_id, occurred_at, gloss_qty, poster_qty, gloss_qty + poster_qty as non_standard_qty
from orders
limit 5;
```

account_id	occurred_at	gloss_qty	poster_qty	non_standard_qty
1001	2015-10-06 17:31:14	22	24	46
1001	2015-11-05 03:34:33	41	57	98
1001	2015-12-04 04:21:55	47	0	47
1001	2016-01-02 01:18:24	32	0	32
1001	2016-02-01 19:27:27	29	28	57

14. **YOUR TURN** Find the percentatge of standard papers ordered for each order. Limit the results to the first 5 orders, and include the id and `account_id` fields.

```
select id, account_id, standard_qty, total, standard_qty/total::float *100 as standard_percent
from orders
limit 5;
```

id	account_id	standard_qty	total	standard_percent
1	1001	123	169	72.78107
2	1001	190	288	65.97222
3	1001	85	132	64.39394
4	1001	144	176	81.81818
5	1001	108	165	65.45455

## Logical Operators

In this section, you will be learning about Logical Operators. Logical Operators include:

**LIKE** This allows you to perform operations similar to using WHERE and =, but for cases when you might not know exactly what you are looking for.

**IN** This allows you to perform operations similar to using WHERE and =, but for more than one condition.

**NOT** This is used with IN and LIKE to select all of the rows NOT LIKE or NOT IN a certain condition.

**AND & BETWEEN** These allow you to combine operations where all combined conditions must be true.

**OR** This allow you to combine operations where at least one of the combined conditions must be true.

15. Find the website for the Whole Foods account.

```
select name, website
from accounts
where website LIKE '%food%';
```

name	website
Tyson Foods	www.tysonfoods.com
US Foods Holding	www.usfoods.com
ConAgra Foods	www.conagrafoods.com
Whole Foods Market	www.wholefoodsmarket.com
Hormel Foods	www.hormelfoods.com
Dean Foods	www.deanfoods.com

Suppose that you are trying to find the website for Whole Foods, and you don't remember their exact name which is (Whole Foods Market. In this case, we can identify all accounts that has the word **food** in their url to find the exact URL for Whole Foods.

The LIKE operator is extremely useful for working with text. You will use LIKE within a WHERE clause. The LIKE operator is frequently used with %. The % tells us that we might want any number of characters leading up to a particular set of characters or following a certain set of characters, as we saw with the **food** syntax above.

16. **YOUR TURN** Find all companies whose names contain the string 'one' somewhere in the name.

```
select name
from accounts
where name LIKE '%one%'
limit 5;
```

name
Honeywell International
INTL FCStone
AutoZone

17. **YOUR TURN** Use the accounts table to find all the companies whose names start with 'C'.

```
select name
from accounts
where name LIKE 'C%'
limit 5;
```

name
CVS Health
Chevron
Costco
Cardinal Health
Citigroup

18. Find the account ID for both Apple and Walmart.

```
select name, id
from accounts
where name in ('Apple', 'Walmart');
```

name	id
Walmart	1001
Apple	1021

The IN operator is useful for working with both numeric and text columns. This operator allows you to use an =, but for more than one item of that particular column. We can check one, two or many column values for which we want to pull data, but all within the same query.

19. **YOUR TURN** Get all account IDs for those companies who were contacted via ads displayed on twitter or via google adwords.

```
select account_id, channel
from web_events
where channel in ('twitter', 'adwords')
limit 5;
```

account_id	channel
1001	adwords
1001	adwords
1001	adwords
1001	twitter
1001	adwords

20. Find the account name, primary poc, and sales rep id for all stores except Walmart, Target, and Nordstrom.

```
select name, primary_poc, sales_rep_id
from accounts
where name not in ('Walmart', 'Target', 'Nordstrom')
limit 5;
```

name	primary_poc	sales_rep_id
Exxon Mobil	Sung Shields	321510
Apple	Jodee Lupo	321520
Berkshire Hathaway	Serafina Banda	321530
McKesson	Angeles Crusoe	321540
UnitedHealth Group	Savanna Gayman	321550

The NOT operator is an extremely useful operator for working with the previous two operators we introduced: IN and LIKE. By specifying NOT LIKE or NOT IN, we can grab all of the rows that do not meet a particular criteria.

21. **YOUR TURN** Find all company account ids who were contacted via any method except using twitter or adwords methods.

```
select account_id, channel
from web_events
```

```
where channel not in ('twitter', 'adwords')
limit 5;
```

account_id	channel
1001	direct
1001	direct
1001	direct
1001	direct
1001	direct

22. **YOUR TURN** Find all the companies whose names do not start with 'C'.

```
select *
from accounts
where name not like 'C%'
limit 5;
```

id	name	website	lat	long	primary_poc	sales_rep_id
1001	Walmart	www.walmart.com	40.23850	-75.10330	Tamara Tuma	321500
1011	Exxon Mobil	www.exxonmobil.com	41.16916	-73.84937	Sung Shields	321510
1021	Apple	www.apple.com	42.29049	-76.08401	Jodee Lupo	321520
1031	Berkshire Hathaway	www.berkshirehathaway.com	40.94902	-75.76390	Serafina Banda	321530
1041	McKesson	www.mckesson.com	42.21709	-75.28500	Angeles Crusoe	321540

23. Pull all orders that occurred between April 1, 2016 and September 1, 2016.

```
select *
from orders
where occurred_at >= '04-01-2016' AND occurred_at <= '09-01-2016'
limit 5;
```

id	account_id	occurred_at	standard_qty	gloss_qty	poster_qty	total	standard_amt_usd	gloss_amt_usd	poster_amt_usd	total_amt_usd
7	1001	2016-04-01 11:20:18	101	33	92	226	503.99	247.17	747.04	1498.20
8	1001	2016-05-01 15:55:51	95	47	151	293	474.05	352.03	1226.12	2052.20
9	1001	2016-05-31 21:22:48	91	16	22	129	454.09	119.84	178.64	752.57
10	1001	2016-06-30 12:32:05	94	46	8	148	469.06	344.54	64.96	878.56
11	1001	2016-07-30 03:26:30	101	36	0	137	503.99	269.64	0.00	773.63

The AND operator is used within a WHERE statement to consider more than one logical clause at a time. Each time you link a new statement with an AND, you will need to specify the column you are interested in looking at. You may link as many statements as you would like to consider at the same time. This operator works with all of the operations we have seen so far including arithmetic operators (+, \*, -, /). LIKE, IN, and NOT logic can also be linked together using the AND operator.

The above query could be written using the BETWEEN operator as follows:

```
select *
from orders
where occurred_at BETWEEN '04-01-2016' AND '09-01-2016'
limit 5;
```

id	account_id	occurred_at	standard_qty	gloss_qty	poster_qty	total	standard_amt_usd	gloss_amt_usd	poster_amt_usd	total_amt_usd
7	1001	2016-04-01 11:20:18	101	33	92	226	503.99	247.17	747.04	1498.20
8	1001	2016-05-01 15:55:51	95	47	151	293	474.05	352.03	1226.12	2052.20
9	1001	2016-05-31 21:22:48	91	16	22	129	454.09	119.84	178.64	752.57
10	1001	2016-06-30 12:32:05	94	46	8	148	469.06	344.54	64.96	878.56
11	1001	2016-07-30 03:26:30	101	36	0	137	503.99	269.64	0.00	773.63

24. **YOUR TURN** Write a query that returns all the orders where the standard\_qty is over 1000, the poster\_qty is 0, and the gloss\_qty is 0.

```
select *
from orders
where standard_qty >1000 AND poster_qty = 0 AND gloss_qty = 0
```

id	account_id	occurred_at	standard_qty	gloss_qty	poster_qty	total	standard_amt_usd	gloss_amt_usd	poster_amt_usd	total_amt_usd
2613	2951	2016-08-15 00:06:12	1171	0	0	1171	5843.29	0	0	5843.29
3260	3491	2014-08-29 22:43:00	1552	0	0	1552	7744.48	0	0	7744.48

25. **YOUR TURN** Which companies who were contacted via twitter or adwords channel and started their account at any point in 2016 sorted from newest to oldest.

```
select *
from web_events
where channel in ('twitter', 'adwords') and occurred_at between '01-01-2016' and '01-01-2017'
order by occurred_at DESC
limit 5;
```

id	account_id	occurred_at	channel
7447	3101	2016-12-31 23:05:47	twitter
8101	3731	2016-12-31 16:03:21	twitter
7689	3331	2016-12-31 07:32:27	twitter
5562	1791	2016-12-31 02:08:50	adwords
7703	3351	2016-12-30 21:06:53	adwords

Using BETWEEN is tricky for dates! While BETWEEN is generally inclusive of endpoints, it assumes the time is at 00:00:00 (i.e. midnight) for dates. This is the reason why we set the right-side endpoint of the period at '01-01-2017'.

26. Find all customers whose orders did not have at least one type of papers.

```
select *
from orders
where standard_qty = 0 OR gloss_qty = 0 OR poster_qty = 0
limit 5;
```

id	account_id	occurred_at	standard_qty	gloss_qty	poster_qty	total	standard_amt_usd	gloss_amt_usd	poster_amt_usd	total_amt_usd
3	1001	2015-12-04 04:21:55	85	47	0	132	424.15	352.03	0	776.18
4	1001	2016-01-02 01:18:24	144	32	0	176	718.56	239.68	0	958.24
11	1001	2016-07-30 03:26:30	101	36	0	137	503.99	269.64	0	773.63
17	1011	2016-12-21 10:59:34	527	14	0	541	2629.73	104.86	0	2734.59
18	1021	2015-10-12 02:21:56	516	23	0	539	2574.84	172.27	0	2747.11

When combining multiple of these operations (AND, OR, BETWEEN, NOT), we frequently might need to use parentheses to assure that logic we want to perform is being executed correctly.

27. Find all customers whose orders did not have at least one type of paper and the order occurred after September 1, 2016. Sort the result from older transactions to the newest.

```
select *
from orders
where (standard_qty = 0 OR gloss_qty = 0 OR poster_qty = 0) AND
      occurred_at >= '2016-09-01'
order by occurred_at
limit 5;
```

id	account_id	occurred_at	standard_qty	gloss_qty	poster_qty	total	standard_amt_usd	gloss_amt_usd	poster_amt_usd	total_amt_usd
6367	3641	2016-09-01 07:48:59	0	134	49	183	0.00	1003.66	397.88	1401.54
6533	4011	2016-09-01 14:28:48	0	0	13	13	0.00	0.00	105.56	105.56
3708	4011	2016-09-01 14:46:17	459	0	0	459	2290.41	0.00	0.00	2290.41
1079	1741	2016-09-02 06:41:11	503	0	34	537	2509.97	0.00	276.08	2786.05
4971	1741	2016-09-02 06:48:59	0	37	0	37	0.00	277.13	0.00	277.13

28. **YOUR TURN** Find list of order ids where either gloss\_qty or poster\_qty is greater than 4000. Only include the id field in the resulting table.

```
select id
from orders
where gloss_qty > 4000 or poster_qty > 4000
limit 5;
```

id
362
731
1191
1913
1939

29. **YOUR TURN** Write a query that returns a list of orders where the standard\_qty is zero and either the gloss\_qty or poster\_qty is over 1000.

```
select id, standard_qty, gloss_qty, poster_qty
from orders
where standard_qty = 0 AND
      (gloss_qty > 4000 or poster_qty > 4000)
```

id	standard_qty	gloss_qty	poster_qty
1913	0	6450	45
4698	0	484	4901
4942	0	10744	95
5791	0	10	11691

30. **YOUR TURN** Find all the company names that start with a 'C' or 'W', and the primary contact contains 'ana' or 'Ana', but it doesn't contain 'eana'.

```
select name, primary_poc
from accounts
where (name LIKE 'C%' OR name LIKE 'W%') AND
      ((primary_poc LIKE '%ana%' or primary_poc LIKE '%Ana%') AND
       primary_poc NOT LIKE '%eana%')
```

name	primary_poc
CVS Health	Anabel Haskell
Comcast	Shana Sanborn